

BCB 726: *Machine Learning for Computational Biology*



The Three Cultures of Data

October 22nd, 2025

“Cultures”?

Statistical Modeling: The Two Cultures


Leo Breiman

Abstract. There are two cultures in the use of statistical modeling to reach conclusions from data. One assumes that the data are generated by a given stochastic data model. The other uses algorithmic models and treats the data mechanism as unknown. The statistical community has been committed to the almost exclusive use of data models. This commitment has led to irrelevant theory, questionable conclusions, and has kept statisticians from working on a large range of interesting current problems. Algorithmic modeling, both in theory and practice, has developed rapidly in fields outside statistics. It can be used both on large complex data sets and as a more accurate and informative alternative to data modeling on smaller data sets. If our goal as a field is to use data to solve problems, then we need to move away from exclusive dependence on data models and adopt a more diverse set of tools.

The Data Modeling Culture

The analysis in this culture starts with assuming a stochastic data model for the inside of the black box. For example, a common data model is that data are generated by independent draws from

response variables = $f(\text{predictor variables, random noise, parameters})$




Model validation. Yes–no using goodness-of-fit tests and residual examination.

Estimated culture population. 98% of all statisticians.

The Algorithmic Modeling Culture

The analysis in this culture considers the inside of the box complex and unknown. Their approach is to find a function $f(\mathbf{x})$ —an algorithm that operates on \mathbf{x} to predict the responses \mathbf{y} . Their black box looks like this:



Model validation. Measured by predictive accuracy.

Estimated culture population. 2% of statisticians, many in other fields.

The Three Cultures of Data

- Statistics
- Machine Learning
- Deep Learning

...and a lot of other ones we won't talk about:

*Bayesian Statistics, Statistical Learning,
Econometrics, Cybernetics, Control Theory, Signal
Processing, Inverse Problems, Operations
Research, Causal Inference, Actuarial Science, &c*

What do I mean by culture?

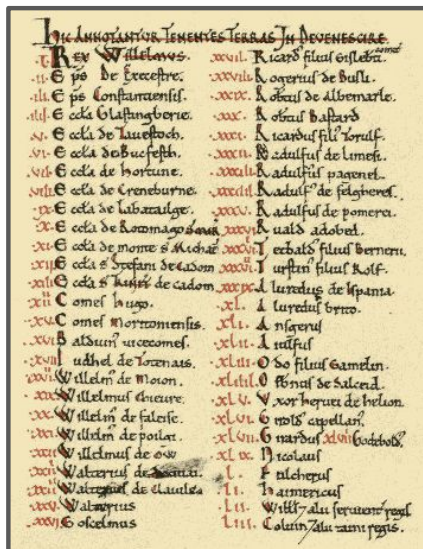
A community of researchers with socially determined assumptions about:

- What is the goal of distilling data into a mathematical model?
- What are the important properties of a data infused mathematical model?

...who then teach classes and write textbooks, creating a socially constructed aura of “the right way to work with data”

*A historical
interlude*

Measurement & estimation aren't new...



William The Conqueror's
Domesday Book of 1086

"What did I just conquer?"

וידבר יהוה אל־מֹשֶׁה לֵאמֹר:

יהוה spoke to Moses, saying:

כִּי תִשָּׂא אֶת־רֹאשׁ בְּנֵי־יִשְׂרָאֵל לִפְקֹדֵיהֶם וְנָתַנוּ אִישׁ כֶּפֶר נַפְשׁוֹ לַיהוָה
בִּפְקֹד אֹתָם וְלֹא־יְהִי בָהֶם נֶגֶף בִּפְקֹד אֹתָם:

When you take a census of the Israelite men according to their army enrollment, each shall pay יהוה a ransom for himself on being enrolled, that no plague may come upon them through their being enrolled.

וְהָיָה כִּלְיָהֶעֱבֹר עַל־הַפְּקֻדִים מִחֻצֵּית הַשֶּׁקֶל בַּשֶּׁקֶל הַקֹּדֶשׁ עֶשְׂרִים
גֶּרָה הַשֶּׁקֶל מִחֻצֵּית הַתְּרוּמָה לַיהוָה:

This is what everyone who is entered in the records shall pay: a half-shekel by the sanctuary weight—twenty *gerabs* to the shekel—a half-shekel as an offering to יהוה.

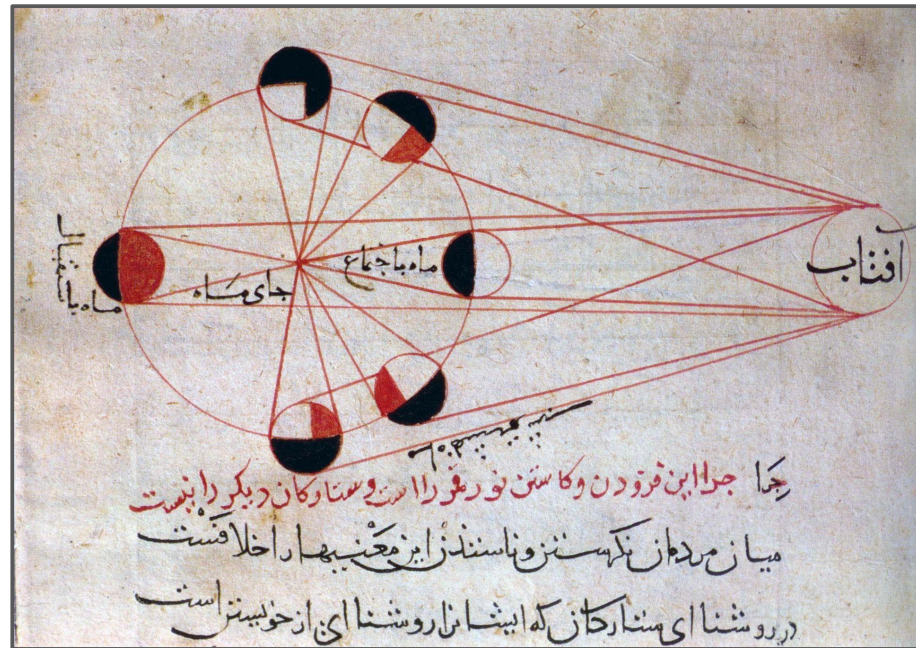
כָּל הָעֹבֵר עַל־הַפְּקֻדִים מִבְּן עֶשְׂרִים שָׁנָה וּמַעְלָה יִתֵּן תְּרוּמַת יְהוָה:

Everyone who is entered in the records, from the age of twenty years up, shall give יהוה's offering:

Beginning of the Torah portion *Ki Tisa*, Moses asked to take a census and collect a tax

Sparks of mathematical modeling of measurement in pre-modern science

- Interplay between **observation** & strong **inductive** priors
- Abstract models but often more descriptive than mathematical
- Validated by agreement with observation but also:
 - Elegance / aesthetics
 - Great “masters”
 - Theology & philosophy



Abu-Rayhan al-Biruni's *Al-Tafhim li Awa'il Sana'at al-Tanjim* (Book on the Elements of Astrology)

Modern science = empiricism

- 1600s science narrowed “natural philosophy” to:
 - **Collect data**
 - **Build mathematical models (repeat)**

“the Universe – which stands continually open to our gaze, but it cannot be understood unless one first learns to comprehend the language and interpret the characters in which it is written. It is written in the language of mathematics” – Galileo in *The Assayer*



Galileo Gallilei's “The Assayer”

Observations of Jupiter's Moons

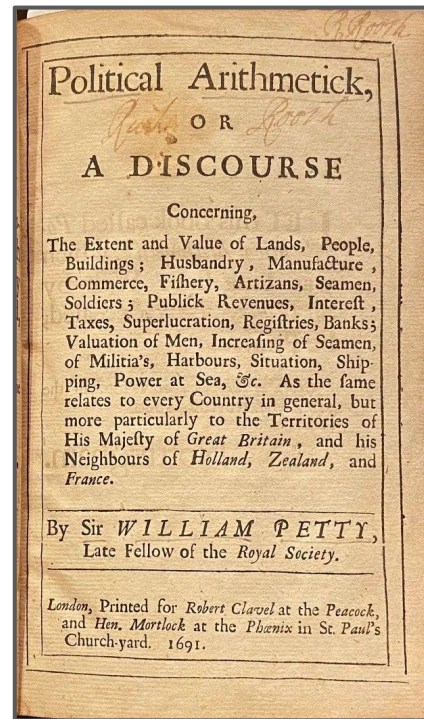
20. Jan. 1610	0 **
30. Jan.	** 0 *
2. Feb.	0 ** *
3. Feb.	0 * *
3. Feb. 5.	* 0 *
4. Feb.	* 0 **
6. Feb.	** 0 *
8. Feb. H. 13.	* * * 0
10. Feb.	* * * 0 *
11.	* * 0 *

Galileo's notebook observing moons of Jupiter (interpreted as evidence for refuting geocentrism)

Statistics = seeing like a state

		The Table of CASUALTIES.																				1860/61		1861/62		1862/63		1863/64		1864/65		1865/66		1866/67		1867/68		1868/69		1869/70		1870/71		1871/72		1872/73		1873/74		1874/75		1875/76		1876/77		1877/78		1878/79		1879/80		1880/81		1881/82		1882/83		1883/84		1884/85		1885/86		1886/87		1887/88		1888/89		1889/90		1890/91		1891/92		1892/93		1893/94		1894/95		1895/96		1896/97		1897/98		1898/99		1899/00		1900/01		1901/02		1902/03		1903/04		1904/05		1905/06		1906/07		1907/08		1908/09		1909/10		1910/11		1911/12		1912/13		1913/14		1914/15		1915/16		1916/17		1917/18		1918/19		1919/20		1920/21		1921/22		1922/23		1923/24		1924/25		1925/26		1926/27		1927/28		1928/29		1929/30		1930/31		1931/32		1932/33		1933/34		1934/35		1935/36		1936/37		1937/38		1938/39		1939/40		1940/41		1941/42		1942/43		1943/44		1944/45		1945/46		1946/47		1947/48		1948/49		1949/50		1950/51		1951/52		1952/53		1953/54		1954/55		1955/56		1956/57		1957/58		1958/59		1959/60		1960/61		1961/62		1962/63		1963/64		1964/65		1965/66		1966/67		1967/68		1968/69		1969/70		1970/71		1971/72		1972/73		1973/74		1974/75		1975/76		1976/77		1977/78		1978/79		1979/80		1980/81		1981/82		1982/83		1983/84		1984/85		1985/86		1986/87		1987/88		1988/89
--	--	--------------------------	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------	--	---------

Table from John Graunt's *Natural and Political Observations Made Upon the Bills of Mortality* (1662)




William Petty's *Political Arithmetic* (1691)

Statistics


So, what is classical statistics?

- There is some real world quantity
 - ...how do we finite noisy measurements into a robust estimate of the “true” value?
- We have a mathematical model of reality
 - ...how do rigorously we use finite noisy measurements to reject (or conditionally accept) the model?

1920s/30s: statistics infiltrates science



CONTENTS		CONTENTS	
viii		ix	
IV. AN AGRICULTURAL EXPERIMENT IN RANDOMISED BLOCKS		50. Interaction of Quantity and Quality 140	
22. Description of the Experiment	55	51. Resolution of Three Comparisons among Four Materials	142
23. Statistical Analysis of the Observations	57	52. An Early Example	143
24. Precision of the Comparisons	64	53. Interpretation of Results	154
25. The Purposes of Replication	66	54. An Experiment with 81 Plots	157
26. Validity of the Estimation of Error	68	IX. THE INCREASE OF PRECISION BY CONCOMITANT MEASUREMENTS. STATISTICAL CONTROL.	
27. Bias of Systematic Arrangements	71	55. Occasions suitable for Concomitant Measurements	167
28. Partial Elimination of Error	72	56. Arbitrary Corrections	173
29. Shape of Blocks and Plots	73	57. Calculation of the Adjustment	176
30. Practical Example	75	58. The Test of Significance	181
V. THE LATIN SQUARE		59. Practical Example	184
31. Randomisation subject to Double Restriction	78	X. THE GENERALISATION OF NULL HYPOTHESES. FIDUCIAL PROBABILITY	
32. The Estimation of Error	81	60. Precision regarded as Amount of Information	187
33. Faulty Treatment of Square Designs	83	61. Multiplicity of Tests of the same Hypothesis	190
34. Systematic Squares	85	62. Extension of the t Test	195
35. Greco-Latin and Higher Squares	90	63. The χ^2 Test	198
36. Practical Exercises	93	64. Wider Tests based on the Analysis of Variance	201
VI. THE FACTORIAL DESIGN IN EXPERIMENTATION		65. Comparisons with Interactions	211
37. The Single Factor	96	XI. THE MEASUREMENT OF AMOUNT OF INFORMATION IN GENERAL	
38. A Simple Factorial Scheme	98	66. Estimation in General	216
39. The Basis of Inductive Inference	106	67. Frequencies of Two Alternatives	218
40. Inclusion of Subsidiary Factors	107	68. Functional Relationships among Parameters	221
41. Experiments without Replication	111	69. The Frequency Ratio in Biological Assay	227
VII. CONFOUNDING		70. Linkage Values inferred from Frequency Ratios	230
42. The Problem of Controlling Heterogeneity	114	71. Linkage Values inferred from the Progeny of Self-fertilised or Intercrossed Heterozygotes	234
43. Example with 8 Treatments, Notation	117	72. Information as to Linkage derived from Human Families	240
44. Design suited to Confounding the Triple Interaction	119	73. The Information elicited by Different Methods of Estimation	244
45. Effect on Analysis of Variance	120	74. The Information lost in the Estimation of Error	247
46. Example with 27 Treatments	123	INDEX 251	
47. Partial Confounding	131		
VIII. SPECIAL CASES OF PARTIAL CONFOUNDING			
48. Dummy Comparisons	138		
49. Dummy Comparisons	138		



CONTENTS	
CHAP.	PAGE
EDITORS' PREFACE	v
AUTHOR'S PREFACE	vii
I. INTRODUCTORY	i
II. DIAGRAMS	27
III. DISTRIBUTIONS	43
IV. TESTS OF GOODNESS OF FIT, INDEPENDENCE AND HOMOGENEITY; WITH TABLE OF χ^2	77
V. TESTS OF SIGNIFICANCE OF MEANS, DIFFERENCES OF MEANS, AND REGRESSION COEFFICIENTS	101
VI. THE CORRELATION COEFFICIENT	138
VII. INTRAClass CORRELATIONS AND THE ANALYSIS OF VARIANCE	176
VIII. FURTHER APPLICATIONS OF THE ANALYSIS OF VARIANCE	211
SOURCES USED FOR DATA AND METHODS	233
INDEX	237
TABLES	
I. AND II. NORMAL DISTRIBUTION	} At End
III. TABLE OF χ^2	
IV. TABLE OF t	
V.A. CORRELATION COEFFICIENT—SIGNIFICANT VALUES	
V.B. CORRELATION COEFFICIENT—TRANSFORMED VALUES	
VI. TABLE OF z	
ix	

Tools of classical statistics

- Experimental design
- Inference & estimators
 - consistent, unbiased, efficient
- Confidence intervals
- Statistical hypothesis testing (Neyman-Pearson)
- Null hypothesis models (Fisher)
- Their unholy marriage: NHST

Biological Question ?



Hypothesis H_0



Design Experiment



Collect Data




Compute p-value



Conclusion

Null Hypothesis Significance Testing



*Teaching Null Hypothesis Significance Testing (NHST) in the Health
Sciences: The Significance of Significance*

Exploratory Data Analysis: You're allowed to look at your data

Exploratory Data Analysis: Past, Present, and Future


John W. Tukey¹

Technical Report No. 302

Princeton University, 408 Fine Hall, Washington Road, Princeton, NJ 08544-1000


Abstract

The 1971-1977 early formulation of Exploratory Data Analysis, in terms of (a) results of some of its techniques and considerations which underlay, at various depths, the choices realized in the books. The 1991-1995 development of Exploratory Analysis of Variance, described in its simplest (two-way table) form and barely sketched in general. Discussion of the changes in apparent philosophy caused by the need to communicate more complicated things, notches, hints, the likely impact on a revised edition of Exploratory Data Analysis 1977. Dreams and targets for what might happen in 1996-2005, with emphasis on Exploratory Regression and the combined use of multiple description.



Exploratory Data Analysis: workflow

- Emphasis on visualization
- Look at residuals of your model
 - ...you can even try different models!
- Does it make sense?
- ...classical stats left for “confirmatory analyses”



Machine Learning

ML: a different lineage

VOL. LIX. No. 236.]

[October, 1950

MIND A QUARTERLY REVIEW OF PSYCHOLOGY AND PHILOSOPHY

I.—COMPUTING MACHINERY AND INTELLIGENCE

By A. M. TURING

1. *The Imitation Game.*

I PROPOSE to consider the question, 'Can machines think?' This should begin with definitions of the meaning of the terms 'machine' and 'think'. The definitions might be framed so as to reflect so far as possible the normal use of the words, but this attitude is dangerous. If the meaning of the words 'machine' and 'think' are to be found by examining how they are commonly used it is difficult to escape the conclusion that the meaning and the answer to the question, 'Can machines think?' is to be sought in a statistical survey such as a Gallup poll. But this is absurd. Instead of attempting such a definition I shall replace the question by another, which is closely related to it and is expressed in relatively unambiguous words.

The new form of the problem can be described in terms of a game which we call the 'imitation game'. It is played with three people, a man (A), a woman (B), and an interrogator (C) who may be of either sex. The interrogator stays in a room apart from the other two. The object of the game for the interrogator is to determine which of the other two is the man and which is the woman. He knows them by labels X and Y, and at the end of the game he says either 'X is A and Y is B' or 'X is B and Y is A'. The interrogator is allowed to put questions to A and B thus:

C: Will X please tell me the length of his or her hair?

Now suppose X is actually A, then A must answer. It is A's

HEURISTIC ASPECTS OF THE ARTIFICIAL INTELLIGENCE PROBLEM*

M. L. Minsky

Introduction

In this report we will discuss, from a heuristic point of view, some of the problems encountered in the design of what might be called "intelligent" machines. We will attempt to indicate, in a few words, the nature of the domain of problems with which we are concerned here.

I do not feel that it would be at all useful to try to lay down an absolute definition of "intelligence" or even of "intelligent behavior". For the things we are trying to accomplish are always related to some set of ad hoc ground rules, problems, and resources. There are certain kinds of performances which, if exhibited by a man, we could all agree embody, or reflect, intelligence. For some purposes we might be able to agree to regard the same performances, in a machine, as "intelligent". But while this convention may be useful in some kinds of discourse, its use in analysis is precluded, for the most part, by two serious faults. First it would constitute a direct evasion of any concise specification of the kinds of activity we are looking for. And then, it seems wrong in spirit, we can often find very simple machines which, for certain tasks, exhibit performances which, if done by a man, we would have to call "intelligent".

Now since we just don't want to confer such a dignity on absurdly simple machines,

* In this paper, which is part of some notes for a book in preparation, the arguments are, for the most part, highly condensed. Details are regularly omitted, particularly those of the mathematical models we have in mind. It is hoped, nevertheless, that it will serve as an introduction to some ideas and techniques that are representative of what is developing in the rapidly growing field of heuristic machines and programs.

ML = symbolic AI isn't working and it seems like AIs need to work with data

- Information Retrieval
- Natural Language Processing
- Computer Vision
- Speech Recognition
- ...and a base of techniques that mix:
 - *Computer Science*
 - *Optimization*
 - *Signal Processing*

A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY

WARREN S. MCCULLOCH and WALTER H. PITTS

Because of the "all-or-none" character of nervous activity, neural events and the relations among them can be treated by means of propositional logic. It is found that the behavior of every net can be described in these terms, with the addition of more complicated logical means for nets containing circles; and that for any logical expression satisfying certain conditions, one can find a net behaving in the fashion it describes. It is shown that many particular choices among possible neurophysiological assumptions are equivalent, in the sense that for every net behaving under one assumption, there exists another net which behaves under the other and gives the same results, although perhaps not in the same time. Various applications of the calculus are discussed.

INTRODUCTION

THEORETICAL neurophysiology rests on certain cardinal assumptions. The nervous system is a net of neurons, each having a soma and an axon. Their adjunctions, or synapses, are always between the axon of one neuron and the soma of another. At any instant a neuron has some threshold, which excitation must exceed to initiate an impulse. This, except for the fact and the time of its occurrence, is determined by the neuron, not by the excitation. From the point of excitation the impulse is propagated to all parts of the neuron. The velocity along the axon varies directly with its diameter, from less than one meter per second in thin axons, which are usually short, to more than 150 meters per second in thick axons, which are usually long. The time for axonal conduction is consequently of little importance in determining the time

Deep Learning: let's imitate the brain to get to AI?

Machine Learning

- Forget models of reality!
- ...let's just do function approximation instead.
- Data is some set of featurized vectors x
 - can have labels (or not)
- Learn functions parameterized by weights
 - minimize a loss function however you can
 - control model complexity to avoid overfitting
- Evaluate predictive accuracy on held-out data

Deep Learning

Like machine learning but...

- terminology inspired by brains
 - “neurons”
 - “layers”
- make the models as big as possible
 - surprisingly, overfitting doesn't matter
- train arbitrary compute graphs with gradient descent

*A Pumpkin
Flavored Example*

Is it a Pumpkin Spice Latte?

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, log_loss, accuracy_score, brier_score_loss

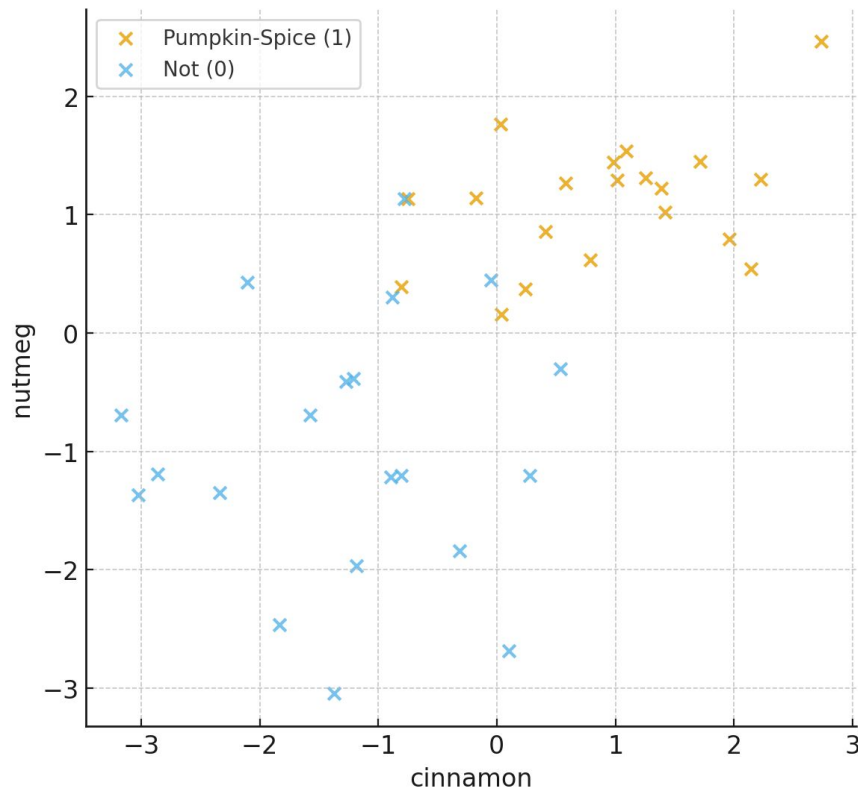
rng = np.random.default_rng(0)

# Tiny, overlapping 2D Gaussians (avoid separation)
n_per = 20
ps_mean = np.array([ 0.8, 0.8]) # Pumpkin-Spice: higher cinnamon/nutmeg
no_mean = np.array([-0.8, -0.8]) # Not
cov = np.array([[1.0, 0.2],
               [0.2, 1.0]])

X_ps = rng.multivariate_normal(ps_mean, cov, size=n_per)
X_no = rng.multivariate_normal(no_mean, cov, size=n_per)
X = np.vstack([X_ps, X_no])
y = np.hstack([np.ones(n_per), np.zeros(n_per)])


df = pd.DataFrame(X, columns=["cinnamon", "nutmeg"]).assign(ps1=y.astype(int))

# Shared train/test split; *all* cultures use the same training data
X_train, X_test, y_train, y_test = train_test_split(
    df[["cinnamon", "nutmeg"]].values, df["ps1"].values,
    test_size=0.4, random_state=0, stratify=df["ps1"].values
)
```



Statistics Flavored Logistic Regression

Dependent variable	Explanatory variables	Model type
Quantitative	1 qualitative with k levels	1-way ANOVA
	Several qualitative	Multi-way ANOVA
	1 or several quantitative	Linear regression
	Mixture of qualitative & quantitative	ANCOVA
Qualitative	1 or several quantitative or qualitative	Logistic regression
Counts with many zeros	1 or several quantitative or qualitative	Log-linear regression



Link: [which modeling method to choose?](#)

```
import numpy as np, pandas as pd, statsmodels.api as sm
from sklearn.metrics import roc_auc_score, log_loss, accuracy_score, brier_score_loss, confusion_matrix

# === Fit ridge (your code) ===
Xtr_sm = sm.add_constant(X_train)
logit = sm.Logit(y_train, Xtr_sm)
res_ridge = logit.fit_regularized(alpha=1.0, L1_wt=0.0) # try alpha in {0.1, 1, 10}

# === Coef table ===
terms = res_ridge.model.exog_names
beta = np.asarray(res_ridge.params)
coef_tbl = pd.DataFrame({"term": terms, "coef": beta, "odds_ratio": np.exp(beta)}).round(3)
print("=== Coefficients (ridge) ===")
print(coef_tbl.to_string(index=False))
```

Inspect everything!

```
# === Coef table ===
terms = res_ridge.model.exog_names
beta = np.asarray(res_ridge.params)
coef_tbl = pd.DataFrame({"term": terms, "coef": beta, "odds_ratio": np.exp(beta)}).round(3)
print("=== Coefficients (ridge) ===")
print(coef_tbl.to_string(index=False))

# === Helper: unpenalized log-likelihood at given beta ===
def loglik_logit(beta, X, y):
    z = X @ beta
    # numerically stable log-sigmoid pieces
    log(sigmoid(z)) = -softplus(-z); log(1-sigmoid(z)) = -softplus(z)
    ll = (y * (-np.log1p(np.exp(-z))) + (1 - y) * (-np.log1p(np.exp(z))))/n
    return float(ll)

# Train diagnostics (unpenalized LL evaluated at penalized beta)
n, k = Xtr_sm.shape
ll_full = loglik_logit(beta, Xtr_sm, y_train)

# Null (intercept-only) model log-likelihood
p0 = y_train.mean()
ll_null = (y_train*np.log(p0) + (1-y_train)*np.log(1-p0)).sum()

deviance = -2.0 * ll_full
mcfadden_r2 = 1.0 - (ll_full / ll_null)
aic = 2*k - 2*ll_full # using unpenalized LL
bic = np.log(n)*k - 2*ll_full

print("\n=== Train diagnostics (evaluated at penalized estimates) ===")
print(f"Log-likelihood (full): {ll_full:.3f}")
print(f"Log-likelihood (null): {ll_null:.3f}")
print(f"Deviance: {deviance:.3f}")
print(f"McFadden R^2: {mcfadden_r2:.3f}")
print(f"AIC (k={k}): {aic:.3f}")
print(f"BIC (k={k}): {bic:.3f}")
print(f"Note: AIC/BIC use the unpenalized log-likelihood at the penalized  $\hat{\beta}$ ; exact penalty-aware ICs require effective df.")

# === Test-set metrics ===
Xte_sm = sm.add_constant(X_test, has_constant='add')
proba = res_ridge.predict(Xte_sm)
pred = (proba >= 0.5).astype(int)

auc = roc_auc_score(y_test, proba)
lloss = log_loss(y_test, proba, labels=[0,1])
acc = accuracy_score(y_test, pred)
brier = brier_score_loss(y_test, proba)
tn, fp, fn, tp = confusion_matrix(y_test, pred).ravel()
sens = tp / (tp + fn) if (tp+fn) else np.nan # recall, TPR
spec = tn / (tn + fp) if (tn+fp) else np.nan # TNR

print("\n=== Test metrics ===")
print(f"AUC: {auc:.3f}")
print(f"LogLoss: {lloss:.3f}")
print(f"Acc: {acc:.3f}")
print(f"Brier: {brier:.3f}")
print(f"Confusion: TP={tp}, FP={fp}, TN={tn}, FN={fn}")
print(f"Sensitivity (TPR): {sens:.3f}")
print(f"Specificity (TNR): {spec:.3f}")
```

```
Optimization terminated successfully (Exit mode 0)
Current function value: 0.2645366576404406
Iterations: 17
Function evaluations: 17
Gradient evaluations: 17
```

=== Coefficients (ridge) ===

	term	coef	odds_ratio
const	0.000	1.000	
x1	0.983	2.673	
x2	2.029	7.606	

=== Train diagnostics (evaluated at penalized estimates) ===

```
Log-likelihood (full): -3.337
Log-likelihood (null): -16.636
Deviance: 6.674
McFadden R^2: 0.799
AIC (k=3): 12.674
BIC (k=3): 16.208
```

Note: AIC/BIC use the unpenalized log-likelihood at the penalized $\hat{\beta}$; exact penalty-aware ICs require effective df.

=== Test metrics ===

```
AUC: 1.000
LogLoss: 0.250
Acc: 0.875
Brier: 0.089
Confusion: TP=8, FP=2, TN=6, FN=0
Sensitivity (TPR): 1.000
Specificity (TNR): 0.750
```

Machine Learning

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression

pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("logreg", LogisticRegression())
])
pipe.fit(X_train, y_train)

pred_probs = pipe.predict_proba(X_test)[:,-1]
pred_labels = (pred_probs >= 0.5).astype(int)

print("AUC:", round(roc_auc_score(y_test, pred_probs), 3))
print("Accuracy:", round(accuracy_score(y_test, pred_labels), 3))
```

AUC: 1.0
Accuracy: 0.812

Deep Learning

```
import torch
from torch import nn
import numpy as np
from sklearn.metrics import roc_auc_score, log_loss, accuracy_score

# Assume X_train, y_train, X_test, y_test are NumPy arrays
torch.manual_seed(0)

# Tensors
Xtr = torch.tensor(X_train, dtype=torch.float32)
ytr = torch.tensor(y_train, dtype=torch.float32).view(-1, 1)
Xte = torch.tensor(X_test, dtype=torch.float32)
yte = torch.tensor(y_test, dtype=torch.float32).view(-1, 1)

# Model: linear -> sigmoid (returns probabilities in (0,1))
model = nn.Sequential(
    nn.Linear(Xtr.shape[1], 1),
    nn.Sigmoid()
)

# Loss takes probabilities because we already applied Sigmoid
loss_fn = nn.BCELoss()
opt = torch.optim.SGD(model.parameters(), lr=0.1)

# Train (full-batch for simplicity)
for _ in range(200):
    p = model(Xtr) # predicted probabilities
    loss = loss_fn(p, ytr)
    opt.zero_grad()
    loss.backward()
    opt.step()

# Evaluate
with torch.no_grad():
    pred_probs = model(Xte).numpy().ravel() # probabilities directly
    pred_labels = (pred_probs >= 0.5).astype(int)

print("AUC:", round(roc_auc_score(y_test, pred_probs), 3))
print("Accuracy:", round(accuracy_score(y_test, pred_labels), 3))
```

```
# Model: linear -> sigmoid (returns probabilities in (0,1))
model = nn.Sequential(
    nn.Linear(Xtr.shape[1], 1),
    nn.Sigmoid()
)
```

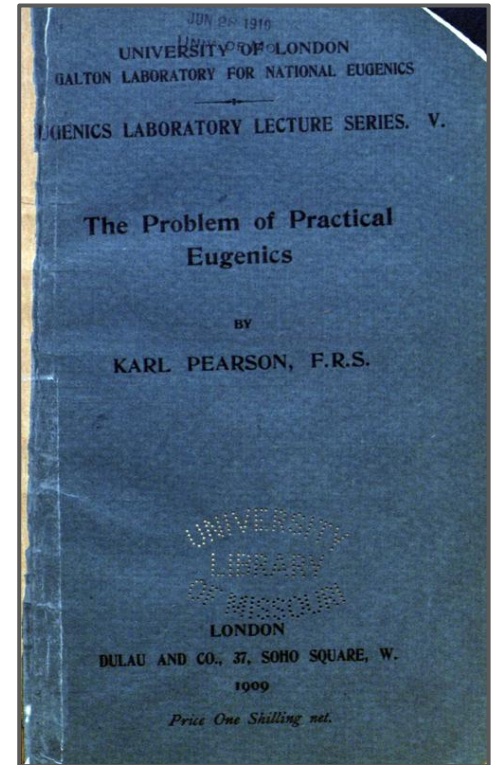
AUC: 1.0
Accuracy: 0.812

❖ *Fin* ❖

Side note: 2/4 “fathers” of statistics were very into eugenics

“...eugenics urges us to simplify our lives, and to simplify our needs; the only luxury worth having is that of a worthy human environment. We must be ready to sacrifice social success, at the call of nobler instincts.” -R. A. Fisher

“History shows me one way, and one way only, in which a high state of civilization has been produced, namely, the struggle of race with race, and the survival of the physically and mentally fitter race.” -Karl Pearson



Bayesian Statistics

- Why are we accepting or rejecting a hypothesis?
- Why do we think there's a single “true” value to parameters?
- Inference should give us a full probability distribution

